

# Flagging terror-related tweets on

Aditi Walia 014516060  
Hardik Kumar 014521039

# INTRODUCTION

- Social media is one of the main means of communication
- Terrorist organisations like ISIS use Twitter to spread their propaganda
- Removing such accounts is a solution but it can't be done manually
- Machine Learning can help in flagging tweets posted by accounts if they are radical

# DATASET

- Data set is scraped from Twitter through GetOldTweets3 and some of it is self-generated
- Against ISIS tweets were found using hashtags like #NoToISIS. ~800 tweets
- Pro ISIS tweets were filtered out from a dataset found online. ~800 tweets
- Random tweets contain random non-radical content ~4000 tweets

Against ISIS

HT\_Felani

I stand with real Iranians, I stand with IR IRAN, I condemn MasihAlinejad for act of terrorism and promoting violence and terrorist activities. #NoToViolence #NoToChaos #NoToTerrorism #NoToPropaganda #NoToFakeNews #NoToGuns #NoToISIS #NoToCyberTerrorism

Pro ISIS

abubakerdimshqi

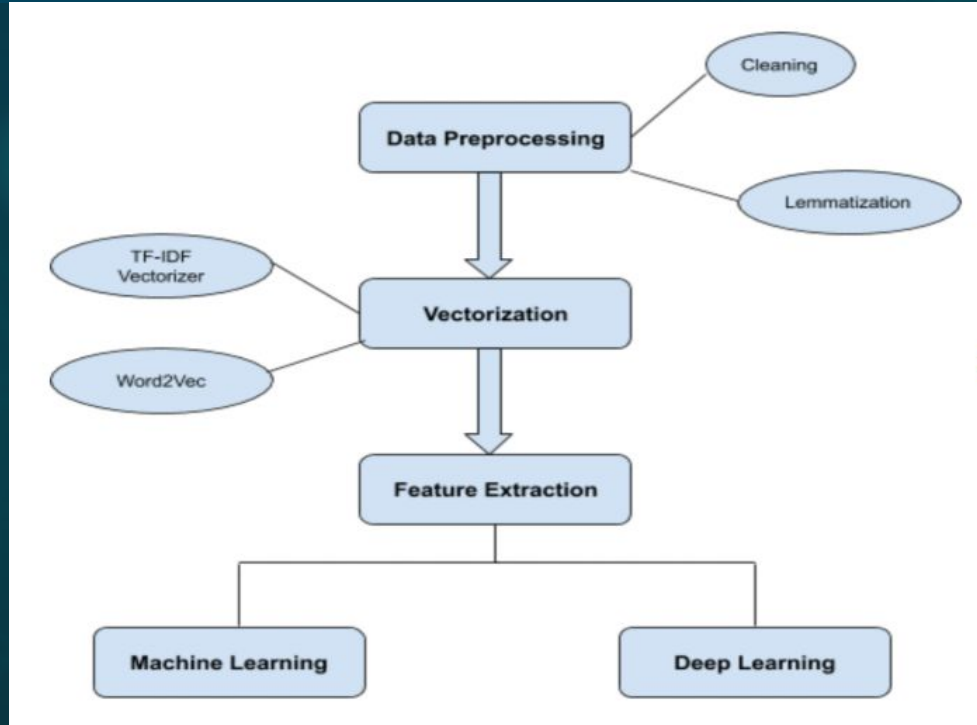
from the Heart love and respect for all #MUJAHIDIN in BILAD #ALSHAM  
Brothers & Sisters you are our hope to lead us to right way #IS

Random

bado0ouri

Should I sleep and skip my classes today

# PROJECT ORGANIZATION



# DATA PREPROCESSING

An important phase; clean the data in order to extract the most useful information out of it

- Remove hyperlinks and “RT” (retweets)
- Remove punctuations
- Remove stop words
- Stemming / Lemmatization

# LEMMATIZING OVER STEMMING

- Reduces inflections or variant forms to base form
- Offers better precision than stemming, because of meaningful chopping of affixes.

Stemming:

```
print(ps.stem('meanness'))  
print(ps.stem('meaning'))|
```

```
mean  
mean
```

Lemmatization:

```
print(wn.lemmatize('meanness'))  
print(wn.lemmatize('meaning'))
```

```
meanness  
meaning
```

# VECTORIZATION

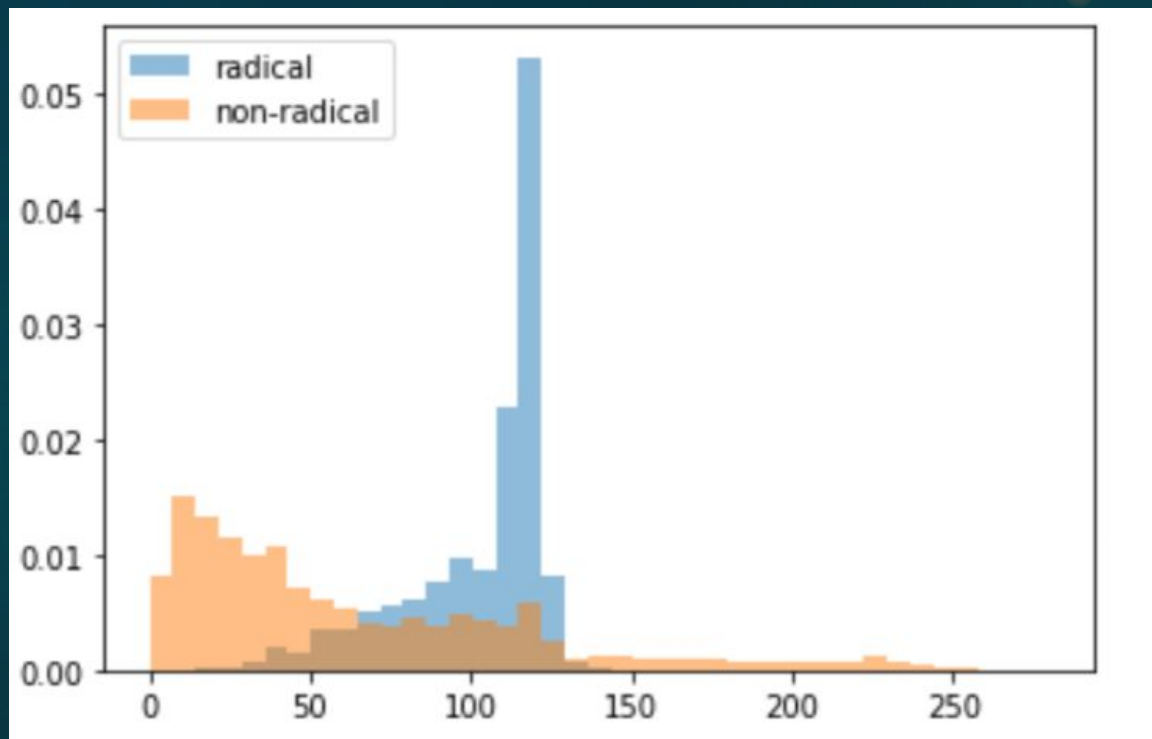
- Word2Vec
- TF-IDF



# FEATURE EXTRACTION

- Tweet length
- Presence of radical keywords
- Word embeddings for Word2Vec vectorization
- TF-IDF feature vectors for TF-IDF vectorization

# TWEET LENGTH VARIATION IN DATASET



# RADICAL KEYWORDS

```
#radical keywords
radical_keywords = ['KUFFUR', 'IS', 'ISLAMIC STATE', '#ILOVEISIS', 'KUFFARS', 'mujahiddeen',
                    'kuffur', 'KUFFAR', 'kuffars', 'kuffar', 'kafir', 'MUJAHIDEEN', 'KUFAR', 'KAFIR',
                    'KUFR', 'mujahideen', '#IS', 'kufar', 'kufr', 'mujahid', 'JIHAD', 'jihad', 'MUJAHID',
                    'MUJAHIDDEEN', '#ISIS', 'Islamic State', '#ILoveISIS', 'ISIL', 'allah', 'Allah', 'Assad', 'assad',
                    'YPG', '#AleppoIsBurning', 'Aleppo', 'martydom', 'Martyrdom']

def radicalWordPresence(text):
    for key in radical_keywords:
        if key in text:
            return 1
    return 0
```

# Word2Vec Embedding

- Embedding words into fixed size vectors
- Common bag of Words
- Skip-n-gram

# Word2Vec EMBEDDING

```
In [12]: model.most_similar('kuffar')
```

```
Out[12]: [('kafir', 0.6600816249847412),  
          ('Kuffar', 0.6480599045753479),  
          ('infidels', 0.6436571478843689),  
          ('unbelievers', 0.6348938941955566),  
          ('polytheists', 0.6023821830749512),  
          ('Kafirs', 0.5989149808883667),  
          ('Allah', 0.5932642817497253),  
          ('infidel', 0.5747469663619995),  
          ('jihad', 0.5737059712409973),  
          ('kafirs', 0.5729705691337585)]
```

# Machine Learning Algorithms on Word2Vec

# Random Forest

Grid Search Results with 5-Fold cross validation on Word2Vec

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_n_estimators	params	split0_test_score	split1_test_score	spl
2	7.117576	0.229126	0.036578	0.004184	None	100	{'max_depth': None, 'n_estimators': 100}	0.911092	0.901408	



# Results

Fit Time: 0.964 / Pred Time: 0.111 ----- Precision: 0.903 / Recall: 0.477 / Accuracy: 0.901

Confusion matrix

```
[[ 93 102]
```

```
 [ 10 931]]
```



PR AUC: 0.735



# Gradient Boosting

Grid Search Results with 5-Fold cross validation on Word2Vec

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_n_estimators	params	split0_test_score	split1_test_score	split
0	129.686026	0.194281	0.021602	0.001401	7	150	{'max_depth': 7, 'n_estimators': 150}	0.931338	0.926937	



# Results

Fit Time: 106.315 / Pred Time: 0.02 ----- Precision: 0.887 / Recall: 0.682 / Accuracy: 0.93

Confusion matrix

```
[[133  62]
```

```
 [ 17 924]]
```



PR AUC: 0.812

# AdaBoost

Grid Search Results with 5-Fold cross validation on Word2Vec

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_learning_rate	param_n_estimators	params	split0_test_score	split1_test_score	s
8	151.377902	7.646242	0.146859	0.014226	1	200	{'learning_rate': 1, 'n_estimators': 200}	0.919894	0.902289	



# Results

Fit Time: 23.471 / Pred Time: 0.11 ----- Precision: 0.848 / Recall: 0.687 / Accuracy: 0.925

Confusion matrix

```
[[134  61]
```

```
 [ 24 917]]
```



PR AUC: 0.794

# K-Nearest Neighbours

Grid Search Results with 5-Fold cross validation on Word2Vec



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_metric	param_n_neighbors	param_weights	params	split0_test_score	split1_test
15	0.097535	0.006593	0.088721	0.009170	manhattan	19	distance	{'metric': 'manhattan', 'n_neighbors': 19, 'weights': 'distance'}	0.880282	0.

# Results

Fit Time: 0.11 / Pred Time: 0.126 ----- Precision: 0.682 / Recall: 0.595 / Accuracy: 0.883

Confusion matrix

```
[[116  79]
 [ 54 887]]
```



PR AUC: 0.673

# Support Vector Machine (SVM)

Grid Search Results with 5-Fold cross validation on Word2Vec

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_kernel	params	split0_test_score	split1_test_score	split2_test_score	split3_t
8	4.519594	0.168714	0.631069	0.014725	1000	rbf	{'C': 1000, 'kernel': 'rbf'}	0.872359	0.857394	0.870599	



# Results

Fit Time: 3.986 / Pred Time: 0.671 ----- Precision: 0.454 / Recall: 0.913 / Accuracy: 0.797

Confusion matrix

```
[[178  17]
```

```
[214 727]]
```





PR AUC: 0.691



# Naive Bayes

## Results

Fit Time: 0.041 / Pred Time: 0.018 ----- Precision: 0.557 / Recall: 0.703 / Accuracy: 0.853  
Confusion matrix  
[[137 58]  
[109 832]]



PR AUC: 0.655

# Summary of Results

Classifier	Recall	Precision	Accuracy
Random Forest	0.477	0.903	0.901
AdaBoost	0.687	0.848	0.925
Gradient Boosting	0.682	0.887	0.93
KNN	0.595	0.682	0.883
SVM	0.913	0.454	0.797
Naive Bayes	0.703	0.557	0.853

# TF-IDF (Term Frequency Inverse Document Frequency)

- Importance of word in document
- Weight assigned between 0 and 1 to each word according to occurrence and importance

$$W_{ij} = tf_{ij} * \log ( N/df_i )$$

- Rarer words occurring in a sentence might be more important to the context of the sentence

# Machine Learning Algorithms on TF-IDF

# Random Forest

Grid Search Results with 5-Fold cross validation on TF-IDF



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_n_estimators	params	split0_test_score	split1_test_score	spl
2	38.610945	1.963792	0.424268	0.085037	None	100	{'max_depth': None, 'n_estimators': 100}	0.943662	0.944542	

# Results

Fit Time: 2.916 / Pred Time: 0.201 ----- Precision: 0.946 / Recall: 0.711 / Accuracy: 0.95

Confusion matrix

```
[[123  50]
 [  7 956]]
```



PR AUC: 0.851

# AdaBoost

Grid Search Results with 5-Fold cross validation on TF-IDF



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_learning_rate	param_n_estimators	params	split0_test_score	split1_test_score	s
4	300.254349	4.757440	11.577173	0.351197	0.5	100	{'learning_rate': 0.5, 'n_estimators': 100}	0.961268	0.950704	

# Results

Fit Time: 53.13 / Pred Time: 3.888 ----- Precision: 0.907 / Recall: 0.85 / Accuracy: 0.964

Confusion matrix

```
[[147  26]  
 [ 15 948]]
```



PR AUC: 0.890



# Gradient Boosting

Grid Search Results with 5-Fold cross validation on TF-IDF

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_n_estimators	params	split0_test_score	split1_test_score	split
0	210.895003	1.266217	0.393229	0.10199	7	150	{'max_depth': 7, 'n_estimators': 150}	0.960387	0.957746	


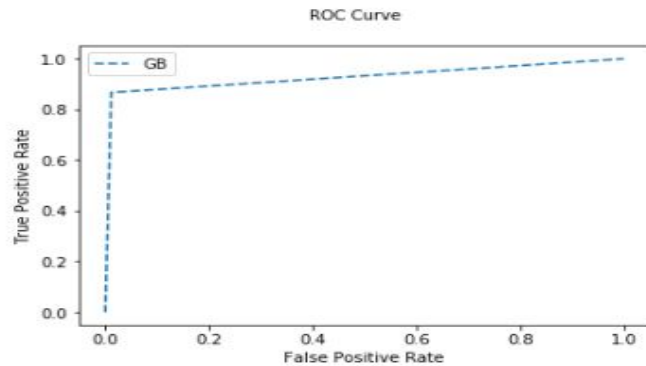
# Results

Fit Time: 348.551 / Pred Time: 0.208 ----- Precision: 0.932 / Recall: 0.867 / Accuracy: 0.97

Confusion matrix

```
[[150  23]
```

```
 [ 11 952]]
```



PR AUC: 0.909

# K-Nearest Neighbours

Grid Search Results with 5-Fold cross validation on TF-IDF


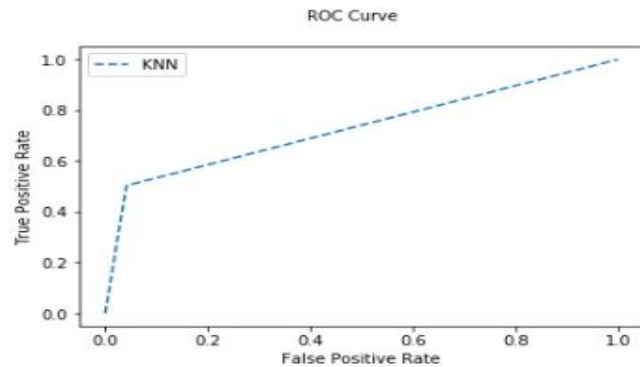
	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_metric	param_n_neighbors	param_weights	params	split0_test_score	split1_test
5	11.454158	1.269725	10.475139	0.636509	euclidean	11	distance	{'metric': 'euclidean', 'n_neighbors': 11, 'weights': 'distance'}	0.889085	0.

# Results

Fit Time: 3.033 / Pred Time: 4.522 ----- Precision: 0.685 / Recall: 0.503 / Accuracy: 0.889

Confusion matrix

```
[[ 87  86]  
 [ 40 923]]
```



PR AUC: 0.632

# Support Vector Machine (SVM)

Grid Search Results with 5-Fold cross validation on TF-IDF

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_kernel	params	split0_test_score	split1_test_score	split2_test_score	split3_t
0	787.271341	121.125993	32.54786	1.464494	100	linear	{'C': 100, 'kernel': 'linear'}	0.955106	0.949824	0.959507	

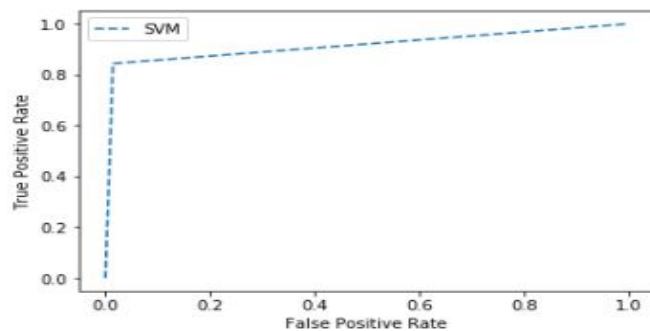
# Results

Fit Time: 365.59 / Pred Time: 26.645 ----- Precision: 0.912 / Recall: 0.844 / Accuracy: 0.964


Confusion matrix

```
[[146 27]
 [ 14 949]]
```

ROC Curve



Precision Recall Curve



PR AUC: 0.890

# Naive Bayes



## Results

Fit Time: 1.862 / Pred Time: 0.404 ----- Precision: 0.38 / Recall: 0.751 / Accuracy: 0.776

Confusion matrix

[[130 43]

[212 751]]



PR AUC: 0.585

# Summary of Results

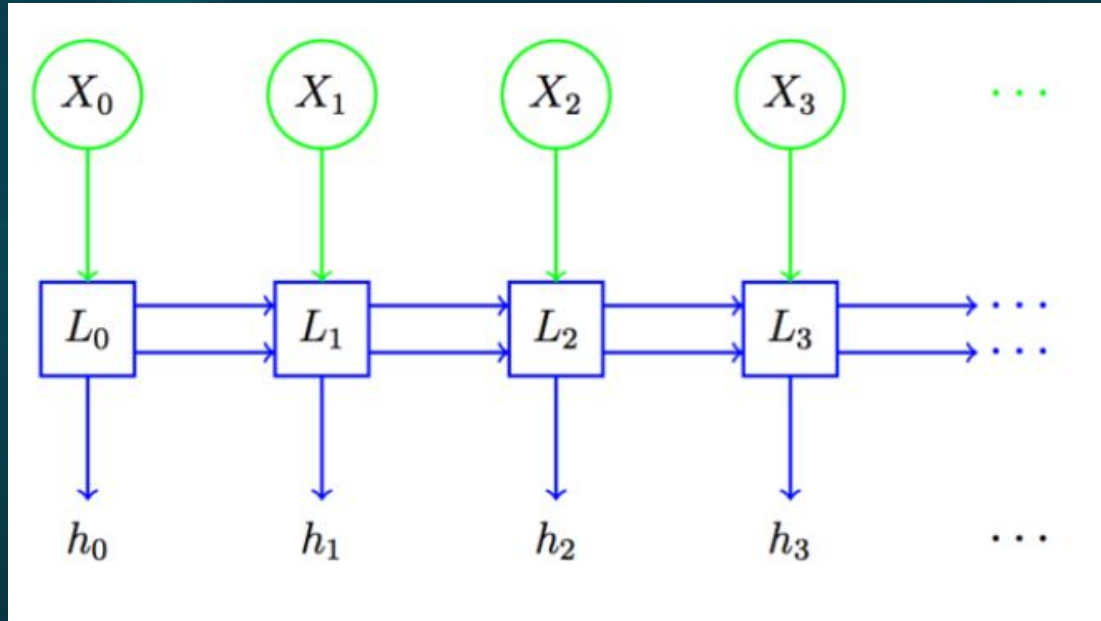
Classifier	Recall	Precision	Accuracy
Random Forest	0.711	0.946	0.950
AdaBoost	0.850	0.907	0.964
Gradient Boosting	0.867	0.932	0.97
KNN	0.503	0.685	0.889
SVM	0.844	0.912	0.964
Naive Bayes	0.751	0.38	0.776



# Deep Learning

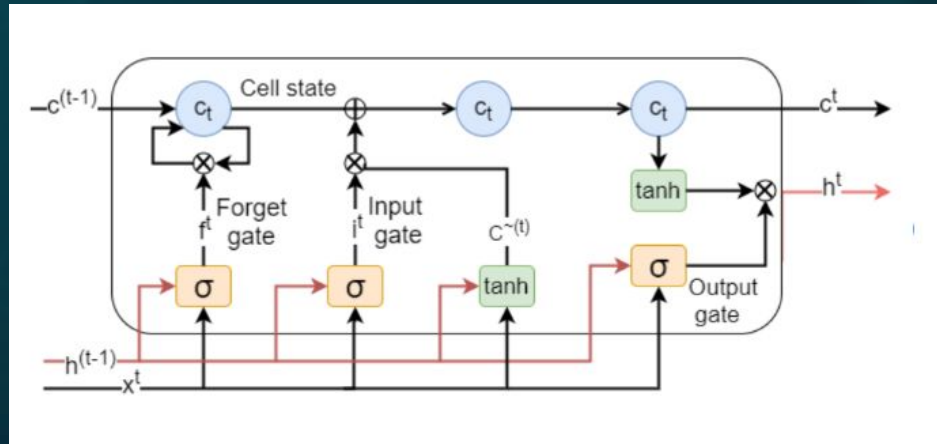
# Long Short Term Memory (LSTM)

Grid Search Results with 5-Fold cross validation on TF-IDF



# Advantages of LSTM

- Mitigates effect of vanishing, exploding, and unstable gradients
- Achieves this through the use of cell state and gates
- Cell state acts as store of long term info, and network decides what to “remember” and what to “forget”



## Results on Word2Vec

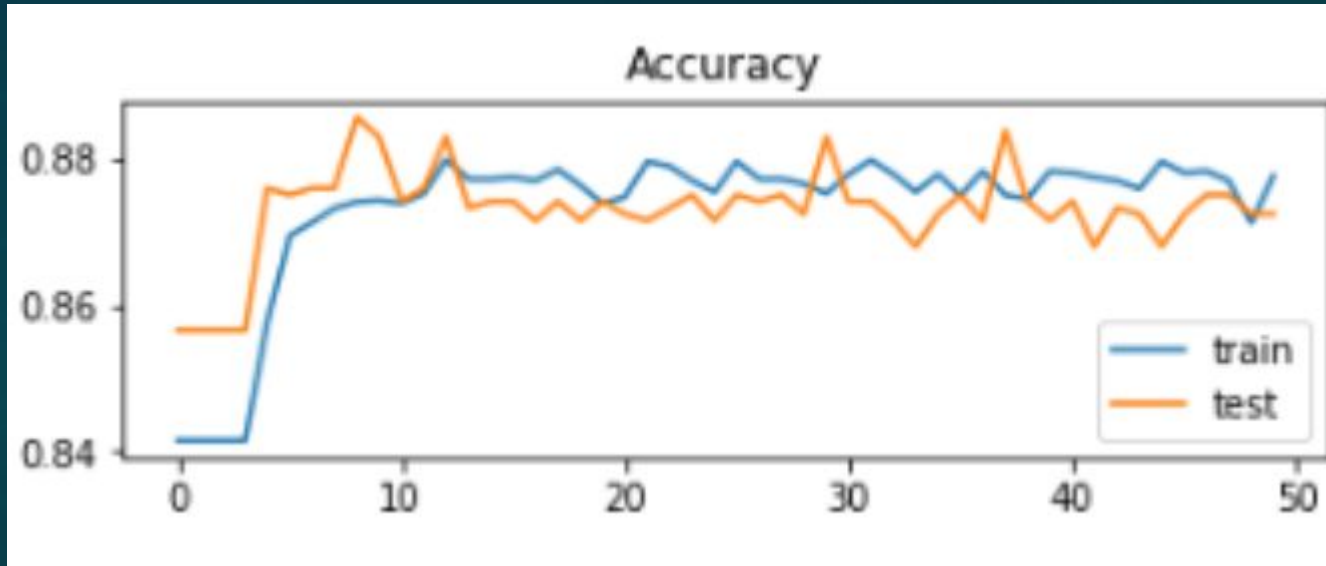
- Used three LSTM layers in the architecture, for 50 epochs
- Focussing on the recall due to the imbalance in data

**Accuracy: 0.872359**

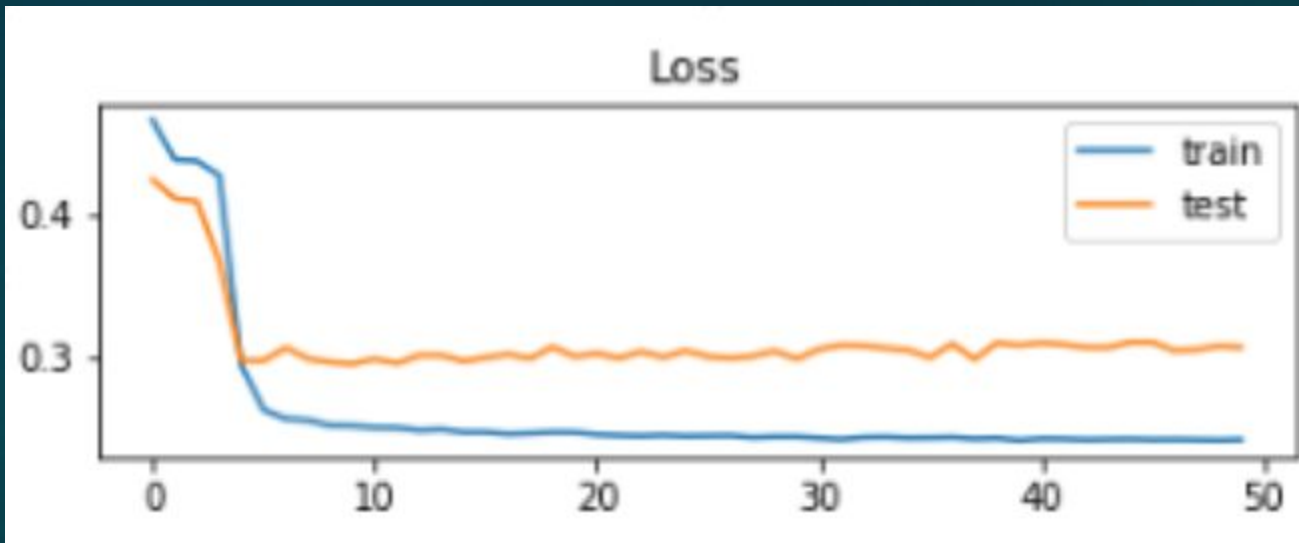
**Precision: 0.554217**

**Recall: 0.564417**

# Accuracy Graph



# Loss Graph



## Results on TFIDF


- Used LSTM layer in the architecture, for 10 epochs
- Focussing on the recall due to the imbalance in data

**Accuracy: 0.882042**

**Precision: 0.595890**

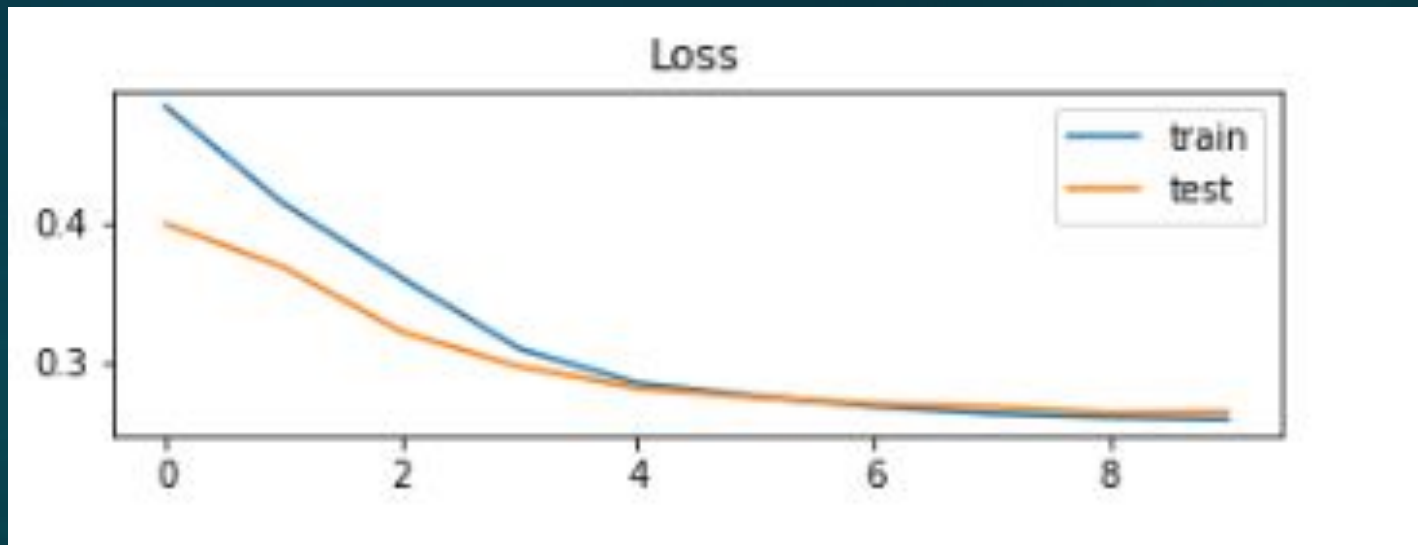
**Recall: 0.537037**

# Accuracy Graph





# Loss Graph



# Conclusions

- For Word2Vec features, **SVM** works the best with a recall of **0.913**
- For TF-IDF features, the boosting algorithms Gradient Boosting and AdaBoost perform the best with recalls of **0.867** and **0.850** respectively. SVM comes close with recall of **0.844**
- SVM performs well because data is linearly separable and because of the abundance of features in our data
- Boosting algorithms performed better due to their iteratively greedy approach of boosting the weak classifiers
- With larger and more diverse dataset, and implementation of deeper architectural-models, we believe that the results might be more accurate
- Also, we used Bag of Words approach in Word2Vec. Using Skip-n-grams might give better results considering data set considers some rarely used words.

